

FlowPools: Lock-Free Deterministic Concurrent Data-Flow Queues

Authors

EPFL, Switzerland
`{firstname.lastname}@epfl.ch`
<http://lamp.epfl.ch>

Abstract. Implementing correct and deterministic parallel programs is hard (additional motivation sentence after this). We present the design and implementation of a fundamental data structure for deterministic parallel data-flow computation. Additionally, we provide a proof of correctness, showing that the implementation is linearizable, lock-free, and deterministic. Finally, we provide microbenchmarks which compare our *flow pools* against corresponding operations on other popular concurrent data structures, in addition to performance benchmarks on a real *XYZ* application using real data. (Keep abstract between 70 and 150 words).

Keywords: data-flow, concurrent data-structure, determinism

1 Introduction

1.1 Motivation

- we want a deterministic model - we do not block in the programming model (i.e. there are no operations which cause blocking until a value becomes available) - we want a non-blocking data-structure (i.e. the operations on the data-structures should themselves be non-blocking) - programs run indefinitely => we need to GC parts of the data structure we no longer need - we want to reduce heap allocation and inline the datastructure as much as possible => lower memory consumption, better cache behaviour and fewer GC cycles

Obligatory multicore motivation paragraph.

Lock-free is better, and why.

Introduction and motivation for data-flow programming model.

2 Model of Computation

Producer-consumer parallelism. Description and image of queue/stream of values, producer, and multiple consumers.

3 Programming Model

The FlowPool suite supports the following operations:

- `Builder.<<(x: T): Builder`
Inserts an element into the underlying FlowPool.
- `Builder.seal(n: Int): Unit`
Seals the underlying FlowPool at `n` elements. The need for the size argument is explained below. A sealed FlowPool may only contain `n` elements. This allows for callback cleanup and termination.
- `FlowPool.doForAll(f: T => Unit): Future[Int]`
Instructs the FlowPool to execute the closure `f` exactly once for each element inserted into the FlowPool (asynchronously). The returned future contains the number of elements in the FlowPool and completes once `f` has been executed for all elements.
- `FlowPool.mappedFold[U, V <: U](acc: V)(cmb: (U,V) => V)(map: T => U): Future[(Int, V)]` Reduces the FlowPool to a single value of type `V`, by first mapping each element to an internal representation using `map` and then aggregating using `cmb`. No guarantee is given about synchronization or order. Returns a future containing a tuple with number of elements in the FlowPool and the aggregated value.
- `FlowPool.builder`
Returns a builder for this FlowPool.
- `Future.map[U](f: T => U)`
Maps this future to another future executing the function `f` exactly once when the first future completes.
- `future[T](f: () => T): Future[T]`
Asynchronously dispatch execution of `f` and return a future with its result.

Determinism of seal We will show that the final size of the FlowPool is required as an argument to the `seal` method in order to satisfy the determinism property of the FlowPool. Look at the following program:

```
val p = new FlowPool[Int]()
val b = p.builder

future {
  for (i <- 1 to 10) { b << i }
  b.seal
}

future { for (i <- 1 to 10) { b << i } }
```

Depending on which for-loop completes first, this program completes successfully or yields an error. A similar program with `b.seal(20)` will always succeed.

Generators In the following we'll present a couple of generators for FlowPools based on common generators in the Scala standard library.

```
def iterate[T](start: T, len: Int)(f: (T) => T) = {
  val p = new FlowPool[T]
  val b = p.builder
  future {
    var e = start
    for (i <- 1 to len) { b << e; e = f(e) }
    b.seal(len)
  }; p
}
```

```
def tabulate[T](n: Int)(f: (Int) => T) = {
  val p = new FlowPool[T]
  val b = p.builder
  future {
    for (i <- 0 to (n-1)) {
      b << f(i)
    }
    b.seal(n)
  }; p
}
```

```
def fill[T](n: Int)(elem: => T) = {
  val p = new FlowPool[T]
  val b = p.builder
  future {
    for (i <- 1 to n) { b << elem }
    b.seal(n)
  }; p
}
```

Monadic Operations In the following we'll present some monadic operations on top of the basic FlowPool operations. This will also show some use-cases of the futures as result type of `doForAll`.

```
def map[S](f: T => S) = {
  val fp = new FlowPool[S]
  val b = fp.builder
  doForAll { x =>
    b << f(x)
  } map { b.seal _ }
  fp
}
```

```

def filter(f: T => Boolean) = {
  val fp = new FlowPool[T]
  val b = fp.builder

  mappedFold(0)(_ + _) { x =>
    if (f(x)) { b << x; 1 } else 0
  } map { case (c,fc) => b.seal(fc) }

  fp
}

def flatMap[S](f: T => FlowPool[S]) = {
  val fp = new FlowPool[S]
  val b = fp.builder

  mappedFold(future(0))(_ <+> _) { x =>
    f(x).doForAll(b << _)
  } map { case (c,cfut) => cfut.map(b.seal _) }

  fp
}

```

where `<+>` is the future of the sum of two `Future[Int]`.

4 Implementation

5 Proofs

5.1 Abstract Pool Semantics

Alex

5.2 Linearizability

Alex

5.3 Lock-Freedom

5.4 Determinism

6 Experimental Results

7 Related Work

Things to probably cite: Oz, gpars, Java CLQ, our futures writeup.

Things we should probably have a look at: Microsoft TPL, Dataflow Java, FlumeJava...

Forcing a bib, [?], [?], [?], [?], [?], [?]

8 Conclusion

Definition 1 (FlowPool). The **FlowPool** is defined as the references *start* and *current*. The **FlowPool state** is defined as the configuration of objects transitively reachable from the reference *start*.

We say that the FlowPool **has** an element e at some time t_0 if and only if the relation $\text{hasElem}(\text{start}, e)$ holds.

We say that a callback f in a FlowPool **will be called** for the element e at some time t_0 if and only if the relation $\text{willBeCalled}(\text{start}, e, f)$ holds.

We say that the FlowPool is **sealed** at the size s at time t_0 if and only if the relation $\text{sealedAt}(\text{start}, s)$ holds.

TODO formal definitions of these relations (involving the flowpool datatypes)

FlowPool operations are **append**, **foreach** and **seal**, and are defined by pseudocodes in figures ...

Definition 2 (Invariants). We define the following invariants for the **Flow-Pool**:

INV1 $\text{start} = b : \text{Block}, b \neq \text{null}, \text{current} \in \text{reachable}(\text{start})$

INV2 $\forall b \in \text{reachable}(\text{start}), b \notin \text{following}(b)$

INV3 $\forall b \in \text{reachable}(\text{start}), b \neq \text{last}(\text{start}) \Rightarrow \text{size}(b) = \text{LASTELEMPOS} \wedge b.\text{array}(\text{BLOCKSIZE} - 1) \in \text{Terminal}$

INV4 $\forall b = \text{last}(\text{start}), b.\text{array} = p \cdot t \cdot n$, where:

$p = X^P, t = t_1 \cdot t_2, n = \text{null}^N$

$x \in \text{Elem}, t_1 \in \text{Terminal}, t_2 \in \{\text{null}\} \cup \text{Terminal}$

$P + N + 2 = \text{BLOCKSIZE}$

INV5 $\forall b \in \text{reachable}(\text{start}), b.\text{index} > 0 \Rightarrow b.\text{array}(b.\text{index} - 1) \in \text{Elem}$

Definition 3 (Validity). A FlowPool state \mathbb{S} is **valid** if and only if the invariants [INV1-5] hold for that state.

Definition 4 (Abstract pool). An **abstract pool** \mathbb{P} is a function from time t to a tuple of sets $(\text{elems}, \text{callbacks}, \text{seal})$ such that:

$\text{seal} \in \{\emptyset, \{w\}\}, w \in \mathbb{N}$

$\text{callbacks} \subset \{(f : \text{Elem} \Rightarrow \text{Unit}, \text{called})\}$

$\text{called} \subseteq \text{elems} \subseteq \text{Elem}$

We say that an abstract pool \mathbb{P} **is in state** $(\text{elems}, \text{callbacks}, \text{seal})$ at time t if and only if $\mathbb{P}(t) = (\text{elems}, \text{callbacks}, \text{seal})$.

Definition 5 (Abstract pool operations). We say that an abstract pool operation op applied to an abstract pool $P_0 = (\text{elems}_0, \text{callbacks}_0, \text{seal}_0)$ at some time t **changes** the state of the abstract pool to $P = (\text{elems}, \text{callbacks}, \text{seal})$ if $\exists t_0, \forall \tau, t_0 < \tau < t, \mathbb{P}(\tau) = P_0$ and $\mathbb{P}(t) = P$.

Abstract pool operation $\text{foreach}(f)$ changes the state at t_0 from $(\text{elems}, \text{callbacks}, \text{seal})$ to $(\text{elems}, (f, \emptyset) \cup \text{callbacks}, \text{seal})$. Furthermore:

$\exists t_1 > t_0, \forall t_2 > t_1, \mathbb{P}(t_2) = (\text{elems}_2, (f, \text{called}_2) \cup \text{callbacks}_2, \text{seal}_2)$, where:

$\text{called}_2 \subseteq \text{elems}_2 \subseteq \text{elems}$